

Project Development Phase
Model Performance Test

Date	18 November 2022
Team ID	PNT2022TMID28091
Project Name	Efficient Water Quality Analysis & Prediction using Machine Learning
Maximum Marks	10 Marks

Model Performance Testing:

Project team shall fill the following information in model performance testing template

Metrics

Classification Model and Regression Model:

Logistic Regression : 0.79

K Nearest Neighbors : 0.78

Decision Tree : 0.86

Random Forest : 0.89

AdaBoost : 0.88

Bagging Classifier : 0.89

XGBoost : 0.87

In [86]: `from sklearn.metrics import classification_report`

```
y_pred_rf= xgb.predict(X_test)
print(classification_report(y_test, y_pred_rf))
```

	precision	recall	f1-score	support
0	0.93	0.90	0.91	518
1	0.66	0.75	0.70	138
accuracy			0.87	656
macro avg	0.80	0.82	0.81	656
weighted avg	0.87	0.87	0.87	656

In [87]: `from sklearn.metrics import classification_report, precision_score, recall_score, confusion_matrix`
`print(precision_score(y_test, y_pred_rf))`
`print(recall_score(y_test, y_pred_rf))`
0.6645161290322581
0.7463768115942029

In [88]: `print(confusion_matrix(y_test, y_pred_rf))`
[[466 52]
 [35 103]]

In [89]: `import pickle`
`filename = 'xgboost.sav'`
`pickle.dump(xgb, open(filename, 'wb'))`

some time later...

Load the model from disk

`n_estimators=500, random_state=42)`

In [84]: `classifiers = [('Logistic Regression', lr), ('K Nearest Neighbours', knn),`
`('Decision Tree', dt), ('Random Forest', rf), ('AdaBoost', ada),`
`('Bagging Classifier', bagging), ('XGBoost', xgb)]`

In [85]: `from sklearn.metrics import accuracy_score`

`for classifier_name, classifier in classifiers:`

`# Fit clf to the training set`
`classifier.fit(X_train, y_train)`

`# Predict y_pred`
`y_pred = classifier.predict(X_test)`
`accuracy = accuracy_score(y_test, y_pred)`

`# Evaluate clf's accuracy on the test set`
`print('{:s} : {:.2f}'.format(classifier_name, accuracy))`

Logistic Regression : 0.79
K Nearest Neighbours : 0.78
Decision Tree : 0.86
Random Forest : 0.89
AdaBoost : 0.88
Bagging Classifier : 0.89
XGBoost : 0.87

Tune the Model

Validation Method –

```
In [82]: grid_knn.fit(X_train, y_train)
         grid_dt.fit(X_train, y_train)
         grid_rf.fit(X_train, y_train)
         grid_ada.fit(X_train, y_train)
         rs_xgb.fit(X_train, y_train)

         print("Best parameters for KNN:", grid_knn.best_params_)
         print("Best parameters for Decision Tree:", grid_dt.best_params_)
         print("Best parameters for Random Forest:", grid_rf.best_params_)
         print("Best parameters for AdaBoost:", grid_ada.best_params_)
         print("Best parameters for XGBoost:", rs_xgb.best_params_)

C:\Users\PC\anaconda3\lib\site-packages\xgboost\sklearn.py:1421: UserWarning: `use_label_encoder` is deprecated in 1.7.0.
  warnings.warn("`use_label_encoder` is deprecated in 1.7.0.")
C:\Users\PC\anaconda3\lib\site-packages\xgboost\sklearn.py:1421: UserWarning: `use_label_encoder` is deprecated in 1.7.0.
  warnings.warn("`use_label_encoder` is deprecated in 1.7.0.")
Best parameters for KNN: {'n_neighbors': 33}
Best parameters for Decision Tree: {'criterion': 'gini', 'max_depth': 21, 'min_samples_leaf': 1}
Best parameters for Random Forest: {'min_samples_leaf': 2, 'n_estimators': 100}
Best parameters for AdaBoost: {'learning_rate': 0.2, 'n_estimators': 400}
Best parameters for XGBoost: {'n_estimators': 1000, 'learning_rate': 0.8}

In [83]: lr = LogisticRegression(random_state=42)
         dt = DecisionTreeClassifier(criterion='gini', max_depth=14, min_samples_leaf=10, random_state=42)
         knn = KNeighborsClassifier(n_neighbors=16)
         rf = RandomForestClassifier(n_estimators=500, min_samples_leaf=2, random_state=42)
         ada = AdaBoostClassifier(n_estimators= 50, learning_rate=0.8)
         xgb = XGBClassifier(n_estimators= 50, learning_rate= 0.5)

         #Let's also apply bagging and boosting
         bagging = BaggingClassifier(DecisionTreeClassifier(criterion='entropy', max_depth=7, min_samples_leaf=2, random_state=42),
                                     n_estimators = 500, random_state = 42)
         bagging.fit(X_train, y_train)
```